

FIELD REPORT ROBINHOOD CHAIN • ID 4663 11 SEP 2026

BUILD WEEK • DAY 6 OF 7

A forecast that can't be edited after the fact

Launch Auditor watches every new token on Robinhood Chain, publishes numbered predictions about how it will go wrong, stamps them on the blockchain before anyone can know the answer, then grades itself in public — alongside the existing scanners. This is the working state, the measurements so far, and the parts that aren't proven yet.

LAUNCHES INDEXED	SIGNED FORECASTS	OUTCOMES GRADED	CELLS CLAIM-READY
1,608 358 live	4,423	963	2 / 11
AGENT BALANCE			
\$46.10			

01

What it is, for someone who has never traded a token

Thousands of new tokens launch every day on this blockchain. Most are worthless within hours. Tools exist that will scan one and tell you it looks risky — but nobody ever checks whether those warnings were right.

Launch Auditor does the checking. For every new token it publishes specific, numbered predictions — *"73% chance the people who created this dump their holdings within 6 hours"* — and writes a fingerprint of that prediction onto the blockchain immediately, where it can't be quietly changed later. Days afterward, it looks at what actually happened and scores itself. Then it publishes the scorecard.

The product isn't the warning. Plenty of tools give warnings. The product is **the track record** — a public, tamper-evident answer to "how often is this thing right?", which no existing scanner offers about itself.

THE SAME THING, PRECISELY

For each launch the agent computes deterministic manipulation and exit-risk features within seconds, publishes separate probabilities for five mechanically-defined outcomes across eleven time horizons, signs the forecast and commits its hash on-chain before the outcome can be known, and later grades every forecast with an open-source scorer against the realized outcome and against public baselines — the base rate, a fixed heuristic, and the existing scanners' own verdicts.

02

Why this gap exists

Token launches used to unfold over days — discovery, then social proof, then a decision to buy. On a chain doing roughly **14,000 launches a day**, all three collapse into the same sixty seconds. Someone posts a contract address; people buy inside a minute; by the time a research page loads the decision is already made.

That reshapes what's scarce. It isn't information about the token — several tools compute that competently. It's **context at the moment of execution**, in a form a machine can consume, from a source with a published error rate.

WHAT ALREADY EXISTS

RobinSight — the closest competitor. Entity clustering by shared funder, holder trim-profiles, local contract audit with evidence. Genuinely good, and more polished than ours on contract analysis.

ScanHood, GoPlus, Robinhood Checker — fast verdict APIs on contract safety and sellability.

WHAT NONE OF THEM DO

- **Commit before the fact.** Every verdict is a live snapshot, editable and unaudited.

- **Publish a hit rate.** No scanner tells you how often it was right.
- **Grade the others.** We score ScanHood and GoPlus as named forecasters in the same table as ourselves.
- **Model a timeline.** They rate the token now; we forecast a specific outcome at a specific hour.

THE HONEST FRAMING

Leading with "we're the benchmark" reads as a ranking claim. The real line is the mechanism: **we publish the forecast before the result, and we grade follower outcomes.** Being the benchmark is a consequence of that, not an assertion.

03

The loop

1 Watch

A
poller
reads
every
Uniswap
pool-
creation
event
on
the
chain.
A
freshness
check
rejects
pools
for

tokens
that
already
existed,
so
pre-
existing
assets
don't
get
scored
as
launches.

2 Measure

Within
seconds:
creator
history,
holder
concentration,
a
wallet
cluster
built
from
on-
chain
evidence,
contract
flags,
v4
hook
permissions
decoded
from
the
hook
address

itself,
and
a
simulated
sell
to
price
the
exit.

3 Forecast & sign

Several
independent
forecasters
produce
probabilities
for
the
same
eleven
outcome
cells.
Each
report
is
canonicalized,
signed
with
the
agent's
key,
and
put
through
a
validator
that
rejects
anything

treating
an
unknown
as
a
pass.

4 Commit

Report
hashes
are
batched
into
a
Merkle
tree
and
the
root
is
posted
to a
registry
contract
on-
chain.
The
forecast's
existence
and
content
are
provable,
and
its
timestamp
precedes
the
outcome

outcome.

5 Grade

At
each
horizon
the
agent
replays
chain
data
to
determine
what
actually
happened,
then
an
open
scorer
computes
calibration
and
discrimination
for
every
forecaster
against
the
same
outcomes.

WHO IS FORECASTING

Five graded participants: `det_v0` (a hand-built statistical model), `det_v0.1` (the same model with intercepts reset to measured base rates), `heuristic_v1` (a deliberately simple rule, as a floor), `base_rate` (the trailing average — the honest baseline any model must beat), and the external scanners mapped onto the same probability scale. A sixth,

11m_deepdive_v0 , is an LLM that investigates qualified launches with read-only chain tools and is scored as just another forecaster — so the dashboard answers whether narrative analysis adds anything over arithmetic.

04

What the measurements say

These are realized base rates from replaying two weeks of history on qualified launches — not predictions, but what actually happened. **The sample sizes are the point.** The scorer refuses to state a claim below 200 graded outcomes and refuses to compute metrics below 100, so most cells are still blank. Showing the blanks is the discipline.

OUTCOME CELL	GRADED	RATE	SHARE OF LAUNCHES	STATUS
TRADING_ALIVE @24h	245	27%		CLAIM-READY
DRAWDOWN_80 @24h	218	34%		CLAIM-READY
LIQ_IMPAIRED @24h	194	41%		6 SHORT
INSIDER_EXIT @6h	151	12%		METRICS ONLY
INSIDER_EXIT @24h	53	11%		TOO THIN
SELL_IMPAIRED @1h / @24h	102	~96%		SUSPECT
All @7d cells · INSIDER_EXIT @72h	0	—	—	NO DATA

Graded = outcomes resolved from chain data with a definite true/false. Cells below 100 can't produce metrics; below 200 the scorer will not let a claim be made. 12,380 more outcomes are queued and resolvable — filling them is compute, not research.

THE FINDING THAT MATTERS

Of qualified launches — ones that cleared a real buyer threshold, not spam — **only 27% are still trading 24 hours later.** Roughly three in four die within a day. One in three loses 80% of its price. Two in five lose most of their liquidity. This is the shape of the thing the product exists to describe, and it's measured, not asserted.

EXISTS TO DESCRIBE, AND IT'S MEASURED, NOT ASSERED.

WHY ONE ROW IS FLAGGED

SELL_IMPAIRED reads ~96% true, which looked implausible. One real bug was found and fixed along the way — the resolver was scoring a failed simulation as a confirmed failure to sell, and a failed call is now *unresolved*, never an outcome. The rate stayed high afterward, so we checked the pools themselves:

LIQUIDITY ON A QUALIFIED LAUNCH, 10 MIN IN	USD
25th percentile	6
Median	1,040
75th percentile	5,267
Share under \$5,000	72%

The test sells a fixed **\$100**. Against a \$1,040 pool that is a tenth of the entire market, which on a concentrated-liquidity pool blows through 30% slippage without difficulty. **96% is therefore real.**

Which converts a suspected bug into a sharper design problem: an outcome that is true 96% of the time carries almost no information. Always predicting "impaired" scores well and tells a trader nothing. The cell either needs the sell sized *relative to pool depth* so it can discriminate, or it should leave the scored set and survive as a descriptive statistic — *a \$100 exit is impaired on 96% of new launches* — which is a striking sentence but not a forecast. **It is excluded from the model until that is decided.**

05

What building it taught us

THE CHAIN IS MOSTLY NOISE

~14,000 launches a day, but only ~8,500 pair against a real quote asset, and of 1,250 sampled historical launches only **161 cleared the qualified bar**. A census of everything is neither affordable nor meaningful; sampling the real ones is.

DECOY POOLS BROKE OUR OWN DATA

44% of sampled launches had a pool with a fee over 10% — junk pools with no trades. Our watcher was recording those as the token's real market, so every downstream price and liquidity check failed. One casualty was a **tokenized NVIDIA stock with \$75M of real liquidity** recorded against a 43%-fee decoy. Now the primary pool is re-derived at report time from the quote asset, fee tier, and actual trading activity.

INFRASTRUCTURE WAS THE REAL CONSTRAINT

Three blockchain data providers in a week. The first had a broken archive; the second throttled hard and mangled large queries; a paid node finally unblocked it. A backfill that appeared to hang for 24 hours turned out to be one unguarded network call with no timeout. Most of the engineering difficulty was here, not in the modeling.

THE INSTRUMENT NEEDS AUDITING TOO

Two separate bugs produced confident, wrong outcomes — a failed call scored as a verdict, and a wrong pool scored as a market. Both were only visible because the pipeline grades itself and the numbers looked implausible. A tool that never checks its own output would have shipped both.

06

The part that isn't about trading

This was built for a contest run by **Orbio**, which grants AI agents inference credits. That imposed a specific test, and it's the most unusual thing in the build: *would this demo work identically on a plain API key? If yes, it proves nothing.*

So the agent manages its own credentials. It mints its own gateway key, polls its balance every 60 seconds, and the size of its daily research budget is a direct function of credits accrued in the last 24 hours — throughput visibly follows the token's activity, and the agent never holds or converts money to make that happen. Every state change is written to a hash-

chained signed lifecycle log. As of this morning it went `NO KEY = ACTIVE` on

chained, signed lifecycle log, as of this morning it went `NO_RETRY` → `ACTIVE` on its own and has been running unattended since.

Registry contract	0xF36F84a7B7DfFB952341d021db51bD76E54fDBEe
Latest batch root	0xe981b60f6c005de6bf1389ee6f79f0b88ddf9f1c26751d96a2395504dfef9720
Agent-minted key	sk-orbio-dY30yq · balance \$46.10 · spendable \$43.10
Artifact commits	12 – weights, feature code, outcome rules, scorer, mappings

The weights, the feature code and the outcome definitions are themselves hashed on-chain, so a change to the rules after a forecast is detectable.

07

What is not proven

Stated plainly, because this is the section an external reviewer should read first.

- **The LLM forecaster has zero grades.** It produces reports that resolve at 24h and 7d. It has not made a live call yet. At submission it will have no track record at all — the newest forecaster in a project about track records.
- **Seven of eleven outcome cells can't support a claim.** Two are ready, two are close, three are thin, four are empty. The fix is compute, not insight, but it isn't done.
- **Unattended continuity is untested overnight.** The credential system works; whether its login token refreshes without a human at a browser is the single biggest open assumption, and the answer arrives tomorrow morning.
- **One outcome definition doesn't discriminate** (`SELL_IMPAIRED`, above). It measures something real but is true 96% of the time, so it can't separate one launch from another. It needs resizing or removal from the scored set.
- **The insider-exit measure is narrow by construction.** It tracks a wallet

cluster built only from direct on-chain evidence. Insiders who funded wallets through paths we don't trace are invisible to it. The number is honest about what it measures and silent about what it misses.

- **Nothing here says the analysis is correct because it is committed.**

Committing proves a forecast existed beforehand. It proves nothing about quality. Quality is what the scorecard is for, and the scorecard is mostly still blank.

WHAT IT EXPLICITLY DOES NOT CLAIM

That it pays for itself. That no human touched the server. That its analysis is right. The three things it does claim: the forecast existed before the outcome, the scorer is reproducible, and the comparison to baselines is public.

08

Where it goes, if it goes anywhere

The interesting direction isn't a better scanner — that market is served. It's the two things this architecture makes possible that a scanner can't do, because both need a commitment mechanism and a wallet index:

A CALLER SCORECARD

Score the people who promote tokens, by outcome. Not follower counts — realized profit and loss for someone who bought a caller's picks at a stated delay. *"Buying this wallet 30 seconds late averaged -12% over the last 20 calls; the wallet itself made +180%."* Feedback from active traders was that this is the number they'd act on, and that it has to be a rolling window because a public edge gets competed away.

AN LP OUTCOME ORACLE

The same event data can compute exactly what a liquidity position at any price range would have earned in fees and lost to impermanent loss. Turned around, that's a forecast — *fees minus IL over the next 24h versus simply holding* — graded like every other forecast. Impermanent loss is one of the

largest quiet losses retail takes, and nobody publishes a scored prediction of it.

Both are roadmap, not built. Neither belongs in contest week.

09

How to pressure-test this

If you want to find the weak points, these are the questions that actually bite. They're ordered by how much damage a bad answer does.

Q1 Who is the user, concretely, and what do they do differently because of this?

The
honest
answer
so
far
is "a
trading
bot
or a
terminal
that
wants
a
demand
gate."
Individual
traders
buying
in
30
seconds
may
not
read
anything.
If

the
answer
is
only
"integrators,"
the
product
is a
B2B
data
feed,
not
a
consumer
tool
—
which
is
fine,
but
changes
everything
downstream.

Q2 Does a published scorecard survive contact with the people it scores?

Callers
and
launchpads
are
the
natural
customers
and
the
natural
subjects.

A
launchpad

whose
pools
score
badly
has
no
incentive
to
surface
the
score.
Terminals
earn
on
volume
whether
trades
win
or
lose.
Who
pays
for a
number
that
reduces
trading?

Q3 Is 245 graded outcomes enough to say anything?

Two
cells
clear
the
project's
own
bar.
Push
on
whether
that

bar
is
high
enough,
and
whether
retrospective
backfill
data
—
replayed
history,
not
live
forecasts
—
should
count
toward
a
track
record
at
all.

Q4 Is the insider-exit cluster too narrow to be meaningful?

12%
of
launches
show
a
creator-
cluster
exit
within
6
hours.
If
real
insiders

route
through
wallets
the
cluster
can't
reach,
that
12%
is a
floor
of
unknown
tightness,
and
the
headline
understates
the
problem
it
claims
to
measure.

Q5 Does on-chain commitment matter to anyone but the builder?

A
trader
wants
"this
usually
dumps
on
you."
Precommitment
is
the
mechanism
that

makes
the
claim
defensible
under
challenge
—
but
if no
one
ever
challenges
it,
it's
expensive